

Project 4: Network Security

This project is split into two parts, with the first checkpoint due on **Wednesday, December 7 at 6:00 PM** and the second checkpoint due on **Friday, December 16 at 6:00 PM**. The first checkpoint is worth 2% of your total grade, and the second checkpoint is worth 10%. We strongly recommend that you get started early.

This is a group project; you **SHOULD** work in **teams of two** and if you are in teams of two, you **MUST** submit one project per team. If two different sets of answers are submitted, the one with lower total grade will be accepted. Please find a partner as soon as possible. If you have trouble forming a team, post on Piazza's partner search forum.

The code and other answers your group submits must be entirely your own work, and you are bound by the Student Code. You **MAY** consult with other students about the conceptualization of the project and the meaning of the questions, but you **MUST NOT** look at any part of someone else's solution or collaborate with anyone outside your group. You **MAY** consult published references, provided that you appropriately cite them (e.g., with program comments), as you would in an academic paper.

Solutions **MUST** be submitted electronically, following the filename and solution formats specified under the submission checklist given at the end of each checkpoint.

"You can't defend. You can't prevent. The only thing you can do is detect and respond."

– Bruce Schneier

Introduction

This project will introduce you to common network protocols, the basics behind analyzing network traces from both offensive and defensive perspectives, and several local network attacks.

Objectives

- Gain exposure to core network protocols and concepts.
- Understand offensive techniques used to attack local network traffic.
- Learn to apply manual and automated traffic analysis to detect security problems.

Guidelines

- You **SHOULD** work in a group of 2.
- Your answers may or may not be the same as your classmates'.
- We generated files for you to submit your answers in. You **MUST** submit your answers in the provided files; we will only grade what's there!
- Each submission file contains an example of expected format. Failure to follow this format may result in 1 point deduction for the section. You **MAY** delete the examples; they will be ignored (along with any other line starting with #) when grading.

Required Tools

- Wireshark 32 bit*
- Python 2.7
- dpkt (Python package)

* We strongly recommend using 32 bit version of Wireshark for Checkpoint 1. Bugs within 64 bit version may prevent you from completing the task.

4.1 Checkpoint 1 (20 points)

Security analysts and attackers both frequently study network traffic to search for vulnerabilities and to characterize network behavior. In this section, you will examine a network trace from a sample network we set up for this assignment. You will search for specific details about the network using Wireshark network packet analyzer (<https://www.wireshark.org>).

You can find your personalized files for this assignment at

[https://cs.princeton.edu/courses/archive/fall16/cos432/hw4/\[netid\]/4.1.1.pcap](https://cs.princeton.edu/courses/archive/fall16/cos432/hw4/[netid]/4.1.1.pcap)

[https://cs.princeton.edu/courses/archive/fall16/cos432/hw4/\[netid\]/4.1.2.pcap](https://cs.princeton.edu/courses/archive/fall16/cos432/hw4/[netid]/4.1.2.pcap), where [netid] is your netid.

4.1.1 Exploring Network Traces (10 points)

Examine the first network trace, 4.1.1.pcap, using Wireshark.

Provide concise answers to the following questions.

1. Identify all the hosts on the local network. (3 points)

- a) What are their IP addresses?
- b) What are their MAC addresses? Assume that each IP address is mapped to one MAC address.

What to submit: Submit a) 4.1.1.1_ip.txt that contains the IP addresses and b) 4.1.1.1_mac.txt that contains the MAC addresses. Write one address per line.

2. How many unique TCP conversations, also known as TCP sessions, are there? (1 points)

What to submit: Submit 4.1.1.2.txt that contains the number of unique TCP conversations. Write the number only.

3. What is the IP address of the gateway? (2 points)

What to submit: Submit 4.1.1.3.txt that contains the IP address of the gateway.

4. Retrieve and submit the content of the files downloaded from the FTP servers. (3 points)

What to submit: Submit a) 4.1.1.4_active.txt that contains the content of the file downloaded from the *active* FTP server and b) 4.1.1.4_passive.txt that contains the content of the file downloaded from the *passive* FTP server.

5. One of the hosts performed port scanning, a technique used to find network hosts that have services listening on one or more target ports.

What is the IP address of the port scanner? (1 points)

What to submit: Submit 4.1.1.5.txt that contains the IP address of the port scanner.

Check your answer formats with the examples provided under *Checkpoint 1: Submission Checklist*.

4.1.2 HTTPS Traffic (10 points)

Examine the second network trace, 4.1.2.pcap, using Wireshark.

Provide concise answers to the following questions.

1. In what year was this traffic captured? (1 points)

What to submit: Submit 4.1.2.1.txt that contains the year. Write the number only.

2. What is the hostname (or Fully Qualified Domain Name) of the server that the client connected to? (1 points)

What to submit: Submit 4.1.2.2.txt that contains the hostname of the website.

3. During TLS handshakes, the client(s) provided a list of supported cipher suites.
List the supported cipher suites from one of the clients. (1 points)

What to submit: Submit 4.1.2.3.txt that contains the list of the supported cipher suites. Write one cipher suite per line. The name and hex Cipher ID of each cipher suite MUST be as shown in Wireshark.

4. Which cipher suite did the server choose for the connection? (1 points)

What to submit: Submit 4.1.2.4.txt that contains the cipher suite chosen by the server. The name and hex Cipher ID of the cipher suite MUST be as shown in Wireshark.

5. A user of the client searched a person's name on the website.
What is the first name of this person? (2 points)

What to submit: Submit 4.1.2.5.txt that contains only the first name of the person searched on the website.

6. The user sent a message to the person found in the previous question.
What is the *body* of the message? (3 points)

What to submit: Submit 4.1.2.6.txt that contains the *body* of the message. The message should be 112 characters long, including whitespaces.

7. Export and submit the user's cookie used in sending the message. (1 points)

What to submit: Submit 4.1.2.7.txt that contains the user's cookie. Your answer should include only the content, not the "Cookie: " prefix.

Check your answer formats with the examples provided under *Checkpoint 1: Submission Checklist*.

Checkpoint 1: Submission Checklist

Make sure that your answers for all tasks up to this point are submitted in the following files by **Wednesday, December 7 at 6:00 PM**.

Team Members

`partners.txt` : a text file containing NETIDs of both members, one NETID per line.
Place the NETID of the student making the submission in the first line.

example content of `partners.txt`

```
your_netid  
partner_netid
```

Solution Format

example content of `4.1.1.1_ip.txt`

```
1.2.3.4  
127.0.0.1
```

example content of `4.1.1.1_mac.txt`

```
0f:0f:0f:0f:0f:0f  
1e:1e:1e:1e:1e:1e
```

example content of `4.1.1.2.txt`

```
461
```

example content of `4.1.1.3.txt`

```
8.8.8.8
```

example content of `4.1.1.4_active.txt`

```
content from active_FTP !@($!@:+_
```

example content of 4.1.1.4_passive.txt

```
content from passive_FTP!@($!@:+_
```

example content of 4.1.1.5.txt

```
192.168.1.254
```

example content of 4.1.2.1.txt

```
1970
```

example content of 4.1.2.2.txt

```
www.example.com
```

```
blog.example.com
```

example content of 4.1.2.3.txt

```
TLS_ECDHE_RSA_WITH_RC4_128_SHA (0xc011)  
TLS_ECDHE_RSA_WITH_3DES_EDE_CBC_SHA (0xc012)  
TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA (0xc013)
```

example content of 4.1.2.4.txt

```
TLS_ECDHE_RSA_WITH_AES_128_CBC_SHA (0xc013)
```

example content of 4.1.2.5.txt

```
john
```

example content of 4.1.2.6.txt

```
=====security is kewl; blah blah blah @#$!)$(@ this message is in a si  
ngle line.=====
```

example content of 4.1.2.7.txt

```
__utma=44258684.1258198627.1456687180.1457965376.1458137415.4; __utmc=442  
58684; __utmz=44258684.1457965376.3.3.utmcsr=iss.illinois.edu|utmccn=(re  
ferral)|utmcmd=referral|utmcct=/  

```

List of files that must be submitted for Checkpoint 1

- `partners.txt`
- `4.1.1.1_mac.txt`
- `4.1.1.1_ip.txt`
- `4.1.1.2.txt`
- `4.1.1.3.txt`
- `4.1.1.4_active.txt`
- `4.1.1.4_passive.txt`
- `4.1.1.5.txt`
- `4.1.2.1.txt`
- `4.1.2.2.txt`
- `4.1.2.3.txt`
- `4.1.2.4.txt`
- `4.1.2.5.txt`
- `4.1.2.6.txt`
- `4.1.2.7.txt`

4.2 Checkpoint 2 (60 points)

4.2.1 Anomaly Detection (40 points)

In this part, you will programmatically analyze trace data to detect port scanning activity.

Port scanning can be used offensively to locate vulnerable systems in preparation for an attack, or defensively for research or network administration. In one port scan technique, known as a SYN scan, the scanner sends TCP SYN packets (the first step in the TCP handshake) and watches for hosts that respond with SYN+ACK packets (the second step).

Since most hosts are not prepared to receive connections on any given port, typically, during a port scan, a much smaller number of hosts will respond with SYN+ACK packets than originally received SYN packets. By observing this effect in a network trace, you can identify source addresses that may be attempting a port scan.

Your task is to develop a Python program that analyzes a PCAP file in order to detect possible SYN scans. You **MUST** use dpkt Python library for packet manipulation and dissection.

For more information about dpkt:

- PyDoc documentation - `pydoc dpkt`
- official website - <https://github.com/kbandla/dpkt>

You may also find this tutorial helpful:

<https://jon.oberheide.org/blog/2008/10/15/dpkt-tutorial-2-parsing-a-pcap-file>

Specifications

- Your program **MUST** take one argument, the name of the PCAP file to be analyzed:
ex) `python2.7 4.2.1.py capture.pcap`
- Your program **MUST** print out the set of IP addresses (one per line) that sent more than 3 times as many SYN packets as the number of SYN+ACK packets they received. For the purpose of this assignment, this rule applies even if the number of packets is very small. For example, following cases are all considered as attacks:

```
SYN=4 ACK+SYN=1
```

```
SYN=4 ACK+SYN=0
```

```
SYN=1 ACK+SYN=0
```

- Each IP address **MUST** be printed only once.
- Your program **MUST** silently ignore packets that are malformed or that are not using Ethernet, IP, and TCP.

Example output: A sample PCAP file captured from a real network can be downloaded at:
https://subversion.ews.illinois.edu/svn/sp16-ece422/_shared/mp4/lbl-internal.20041004-1305.port002.dump.anon.pcap

(Source: <ftp://ftp.bro-ids.org/enterprise-traces/hdr-traces05>)

For this input, your program's output **MUST** be these lines, in any order:

```
128.3.23.2
128.3.23.5
128.3.23.117
128.3.23.158
128.3.164.248
128.3.164.249
```

What to submit: Submit `4.2.1.py`, a Python program that accomplishes the task specified above. You **SHOULD** assume that the grader uses `dpkt 1.8`. You **MAY** use standard Python system libraries, but your program **SHOULD** otherwise be self-contained. We will grade your detector using a variety of different PCAP files.